



Wenhao Wu (吴文浩)

Master's Student | Nanjing University

(+86) 18005527326 Email: wenhaowu@smail.nju.edu.cn

Homepage: wuwenhao09.github.io Google Scholar

Research Interests

Large Language Model Agents, Retrieval-Augmented Generation, Meta Reinforcement Learning, and Vision-Language-Action.

Education

Nanjing University, School of Management and Engineering 09/2024–Present
M.Eng. in Control Engineering; supervised by Prof. Zhi Wang.

Nanjing University, School of Management and Engineering 09/2020–06/2024
B.Eng. in Automation.

Research Experience

- **MemFuse: Multi-Source Memory Fusion from Fragmented Observations**
arXiv preprint, 2026; co-author
Chao Li, Yuanfa Li, **Wenhao Wu**, Xule Liu, Zhi Wang, Kun Shao. Introduces MemFuseBench for multi-source memory fusion and a dual-layer causal fusion graph that preserves source provenance while organizing fragmented observations into retrievable memories.
- **From Conflict to Consensus: Boosting Medical Reasoning via Multi-Round Agentic RAG**
ICML 2026; first author
Wenhao Wu, Zhentao Tang, Yafu Li, Shixiong Kai, Mingxuan Yuan, Zhenhong Sun, Chunlin Chen, Zhi Wang. Proposed conflict-guided multi-round retrieval for medical reasoning, using disagreement among sampled responses to generate precise queries and sequence-level entropy or a lightweight BERT verifier to rerank historical reasoning trajectories.
- **Look Before You Leap: Distilling Tree Search into Action Evaluation for Frozen VLA Models**
arXiv preprint, 2026; co-author
Xinyi Xie, Zican Hu, Zhanyu Liu, Yicheng Dong, **Wenhao Wu**, Haoran Li, Chunlin Chen, Zhi Wang*, Pichao Wang. Proposed SVA (Search, Value, and Act), which keeps the VLA backbone frozen and distills Monte-Carlo tree search into a lightweight Q-value evaluator that selects the best of N candidate actions at deployment without simulator access.
- **Mixture-of-Experts Meets In-Context Reinforcement Learning**
NeurIPS 2025; first author
Wenhao Wu, Fuhong Liu, Haoru Li, Zican Hu, Daoyi Dong, Chunlin Chen, Zhi Wang. Proposed T2MIR, a dual-MoE framework with Token-wise and Task-wise experts for multi-modal trajectory processing and multi-task in-context RL.
- **Text-to-Decision Agent: Offline Meta-Reinforcement Learning from Natural Language Supervision**
NeurIPS 2025; co-first author
Shilin Zhang*, Zican Hu*, **Wenhao Wu***, et al. Aligned trajectory decision embeddings with natural-language supervision through a world-model encoder, contrastive LoRA adaptation, and conditional policies.
- **Meta-DT: Offline Meta-RL as Conditional Sequence Modeling with World Model Disentanglement**
NeurIPS 2024; student second author
Zhi Wang, Li Zhang, **Wenhao Wu**, Yuanheng Zhu, Dongbin Zhao, Chunlin Chen. Introduced a context-aware world model and self-guided/complementary prompts to improve Decision Transformer generalization to new tasks.

Competition Experience

TRACE: A Self-Evolving Skill Bank for Consistent, Limit-Aware LLM Agents

[Technical Report](#) | [Homepage](#) | [Code](#) | [Leaderboard](#)

Darwin Agent · **Wenhao Wu***, Menghao Zhang*, Xin Wang*, Zhi Wang, Kun Shao, Jian Luan.

Technical report for the CAR-bench Challenge Open Track at IJCAI-ECAI 2026. Developed a model-agnostic framework that evolves a modular Skill Bank through trajectory-contrastive analysis and performs state-conditioned skill orchestration. The Darwin Agent team won both the **Rank Award** and **Innovation Award**, ranking **1st in the Open Track** with 70.0% Pass³ on the official hidden set.

Research Internship

Huawei Noah's Ark Lab, Research Intern

07/2025–01/2026

Worked on medical RAG and agentic reasoning: built a medical knowledge retrieval service; designed conflict-guided adaptive retrieval; fine-tuned a BERT-based answer verifier and reranker; explored reflection-based adaptive stopping and Agentic RL for medical reasoning.

Honors and Certifications

- People's Scholarship, Nanjing University (2 times).
- First-Class Academic Scholarship for master's students (2 times).
- High-value scholarship (1 time).

Academic Service

Gold Reviewer, ICML 2026; Reviewer, NeurIPS 2026.